

Optimization Theory

Piotr Więcek,
Faculty of Pure and Applied Mathematics,
Wrocław University of Science and Technology

21.10.2020

Gradient methods: Part I

Gradient methods: general idea I

- We want to find a local unconstrained minimum x^* of a given function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ by constructing a sequence of approximations of x^* , x_0 (given by the user), x_1 , x_2 , x_3, \dots converging to x^* .
- For any n , $x_n = F(x_{n-1})$, where the form of function F should depend on f , ∇f , possibly $\nabla^2 f$.
- **Desired properties:**
 - 1 Each iteration decreases the value of f .
 - 2 Decrease in the value of f big enough not to allow convergence before x^* is reached.

Gradient methods: general idea II

Gradient methods are constructed as follows:

For a given starting point x_0 next iterations are given by the formula

$$x_{k+1} := x_k - \alpha_k D_k \nabla f(x_k),$$

where:

D_k is the $n \times n$ **descent direction matrix** (It allows to choose the direction in which the value of f will decrease)

$\alpha_k \in (0, +\infty)$ is the **stepsize** (It allows to manipulate the distance between x_k and x_{k+1} , so that the method does not stop before reaching the minimizer).

Descent direction selection: general rule

Theorem

Suppose D is symmetric positive definite. Then for any $x \in \mathbb{R}^n$ such that $\nabla f(x) \neq 0$,

$$f(x - \alpha D \nabla f(x)) < f(x)$$

for α small enough.

By choosing each D_k positive definite we guarantee that in each iteration we decrease the value of f .

Descent direction selection: steepest descent I

The simplest positive definite matrix that we know is the **identity matrix** I .

The method where $D_k = I$ for any k is called the **steepest descent method**.

Why is it called the steepest descent?


From mathematical analysis we know that $\nabla f(x)$ gives the direction in which the slope of the function is maximized.

So by taking $x_{k+1} = x_k - \alpha_k \nabla f(x_k)$ we should decrease the value of f the fastest possible!

Descent direction selection: steepest descent II

Why do we need any other methods then?

The steepest descent method is known to **zigzag** like on this picture:



It makes its convergence to the minimum very slow.

Descent direction selection: Newton's method I

Alternative choice of the descent direction is given by the **Newton (Newton-Raphson)** method:

$$D_k := \left(\nabla^2 f(x_k) \right)^{-1} \quad \text{for } k = 0, 1, 2, \dots$$

D_k defined in such a way need not be positive definite!

\implies when implementing the Newton method we should **replace** $(\nabla^2 f(x_k))^{-1}$ **by** I or a convex combination of the two **whenever** the former is not positive definite.

Descent direction selection: Newton's method II

Where does Newton descent direction form come from?

- The Taylor expansion of function f around x_k is:

$$f(x) = f(x_k) + \nabla f(x_k)^T (x - x_k) + \frac{1}{2} (x - x_k)^T \nabla^2 f(x_k) (x - x_k) + o(\|x - x_k\|^2)$$

- If we get rid of the last term, we get the quadratic approximation of f around x_k :

$$f(x) \approx f(x_k) + \nabla f(x_k)^T (x - x_k) + \frac{1}{2} (x - x_k)^T \nabla^2 f(x_k) (x - x_k).$$

- We can now look for the minimum of this polynomial:

$$\nabla f(x_k) + \nabla^2 f(x_k) (x - x_k) = 0 \implies x = x_k - \left(\nabla^2 f(x_k) \right)^{-1} \nabla f(x_k).$$

Newton's method chooses the direction where the minimum of the quadratic approximation of f lies.

There are several ways of choosing the stepsize α_k .

1 Optimal stepsize (minimization rule):

We choose α_k minimizing the function

$$F_k(\alpha) := f(x_k - \alpha D_k \nabla f(x_k)) \text{ over } (0, +\infty)$$

2 Limited optimization rule:

We choose α_k minimizing the function

$$F_k(\alpha) := f(x_k - \alpha D_k \nabla f(x_k)) \text{ over } (0, s] \text{ for some given constant } s.$$

Usually, we are not able to apply these rules using analytical tools, thus we implement them with the aid of **line search algorithms** (optimization algorithms for one-dimensional problems).


Line search algorithms I: interval reduction rules

The idea:

Suppose the function $f : \mathbb{R} \rightarrow \mathbb{R}$ is unimodal on interval $[a, b]$ with a minimum at point x^* that we want to find, and the points x_a, x_b are such that $x_L := a < x_a < x_b < b =: x_U$. Then there are 3 possibilities:

- $f(x_a) < f(x_b)$,
- $f(x_a) > f(x_b)$,
- $f(x_a) = f(x_b)$.

In each of these cases we can shrink the original search interval.



Line search algorithms II: interval reduction rules cont'd

In general, the following rules can be applied:

- $f(x_a) < f(x_b) \implies x^* \in [x_L, x_b]$,
- $f(x_a) > f(x_b) \implies x^* \in [x_a, x_U]$,
- $f(x_a) = f(x_b) \implies x^* \in [x_a, x_b]$.

If we decrease the interval in such a way iteratively, we may shrink it to an arbitrarily small interval containing x^* .

Line search algorithms III: golden-section search I

This can be implemented in the **golden-section search method**:

Input: $a, b, f, \varepsilon > 0$

$x_L := a, x_U := b, K := \frac{\sqrt{5}-1}{2}$

while $x_U - x_L > \varepsilon$ do

$x_a := x_U - K(x_U - x_L), x_b := x_L + K(x_U - x_L),$

$f_a = f(x_a), f_b = f(x_b),$

 if $f_a < f_b$

$x_U := x_b,$

 elseif $f_a > f_b$

$x_L := x_a$

 else

$x_L := x_a, x_U := x_b$

$x^* := \frac{x_U + x_L}{2}$

Line search algorithms IV: golden-section search II

The same algorithm, but with less computations of function f :

```
 $x_L := a, x_U := b, K := \frac{\sqrt{5}-1}{2}$   
 $x_a := x_U - K(x_U - x_L), x_b := x_L + K(x_U - x_L),$   
 $f_a = f(x_a), f_b = f(x_b),$   
while  $x_U - x_L > \varepsilon$  do  
  if  $f_a < f_b$   
     $x_U := x_b,$   
     $x_b := x_a, x_a := x_U - K(x_U - x_L)$   
     $f_b := f_a, f_a := f(x_a)$   
  elseif  $f_a > f_b$   
     $x_L := x_a$   
     $x_a := x_b, x_b := x_L + K(x_U - x_L)$   
     $f_a := f_b, f_b := f(x_b)$   
  else  
     $x_L := x_a, x_U := x_b$   
     $x_a := x_U - K(x_U - x_L), x_b := x_L + K(x_U - x_L),$   
     $f_a := f(x_a), f_b := f(x_b)$   
 $x^* := \frac{x_U + x_L}{2}$ 
```

Line search algorithms V: quadratic interpolation I

The idea:

- For any three points $(a, f(a))$, $(b, f(b))$, $(c, f(c))$ such that $a < c < b$ there is a unique polynomial of 2nd degree p such that $p(a) = f(a)$, $p(b) = f(b)$, $p(c) = f(c)$.
- Assuming that f behaves similarly to p , we may expect that the minimum of f is close to the minimum of p .
- If the approximation is not good enough, we may use the minimizer \bar{x} of polynomial p (together with c) to reduce the search interval.
- We may repeat the same procedure for the new interval.

Line search algorithms VI: quadratic interpolation II

Step 1: We search along the line to find initial points $a < c < b$ such that $f(a) > f(c) < f(b)$.

Step 2:

$x^* := b + \varepsilon, \bar{x} := c$

while $|x^* - \bar{x}| > \varepsilon$ do

$x^* := \bar{x}$

$\bar{x} := \frac{(b^2 - c^2)f(a) + (c^2 - a^2)f(b) + (a^2 - b^2)f(c)}{2[(b - c)f(a) + (c - a)f(b) + (a - b)f(c)]}$

 if $\bar{x} > c$

$d := \bar{x}$

 else

$d := c, c := \bar{x}$

 if $f(c) < f(d)$

$b := d$

 else

$a := c, c := d$

$x^* := \bar{x}$

Line search algorithms VII: cubic interpolation I

The idea:

- For any **two** points $(a, f(a))$, $(b, f(b))$ such that $a < b$ there is a unique polynomial of 3rd degree p satisfying $p(a) = f(a)$, $p(b) = f(b)$, $p'(a) = f'(a)$, $p'(b) = f'(b)$.
- We can approximate the minimum of f by a local minimum of p . A 3rd degree polynomial usually has one local minimum and one local maximum. We want the minimum to be in $[a, b]$. For this to be sure f should satisfy:

$$f'(a) < 0 \text{ and } (f'(b) \geq 0 \text{ or } f(b) \geq f(a)).$$

- If the approximation using the minimizer of p , \bar{x} , is not good enough, we can divide the interval $[a, b]$ into two subintervals $[a, \bar{x}]$ and $[\bar{x}, b]$ and continue the search on one of them.

Line search algorithms VIII: cubic interpolation II

Step 1: We search along the line to find initial points $a < b$ such that $f'(a) < 0$ and $(f'(b) \geq 0$ or $f(b) \geq f(a))$.

Step 2:

$$x^* := b + \varepsilon, \bar{x} := a - \varepsilon$$

while $|x^* - \bar{x}| > \varepsilon$ do

$$x^* = \bar{x}$$

$$z := \frac{3(f(a)-f(b))}{b-a} + f'(a) + f'(b), \quad w = \sqrt{z^2 - f'(a)f'(b)}$$

$$\bar{x} := b - \frac{f'(b)+w-z}{f'(b)-f'(a)+2w}(b-a)$$

if $f'(\bar{x}) \geq 0$ or $f(\bar{x}) \geq f(a)$

$$b := \bar{x}$$

else

$$a := \bar{x}$$

$$x^* := \bar{x}$$

Stepsize selection II

3 Armijo rule (inexact line search)


We take the constants $s > 0$ (the end of the search interval), $\beta \in (0, 1)$ (the rate of reducing the search interval) and $\sigma \in (0, 1)$ (slope reduction factor), and make the loop:

$\alpha := s$

while $f(x_k - \alpha D_k \nabla f(x_k)) \geq f(x_k) - \sigma \alpha \nabla f(x_k)^T D_k \nabla f(x_k)$ do

$\alpha := \beta \alpha$

The final value of α is our α_k .



Stepsize selection III

4 n -dimensional quadratic approximation:

We take the value minimizing the Taylor 2nd order approximation of $f(x_k - \alpha D_k \nabla f(x_k))$:

$$\alpha_k := \frac{\nabla f(x_k)^T \nabla f(x_k)}{\nabla f(x_k)^T \nabla^2 f(x_k) \nabla f(x_k)} \quad \text{for } k = 0, 1, 2, \dots$$

for the steepest descent method or

$$\alpha_k := 1 \quad \text{for } k = 0, 1, 2, \dots$$

for Newton's method.

5 Diminishing stepsize:

Any sequence α_k such that $\alpha_k \searrow 0$ and $\sum_{k=0}^{\infty} \alpha_k = +\infty$.

6 Constant stepsize:

$$\alpha_k \equiv \text{const.}$$